

Java Graphics Tutorial

Case Study on Tic-Tac-Toe Part 2: With AI


TABLE OF CONTENTS (HIDE)

1. Playing Against Computer with AI
 - 1.1 Implementing the AI Player
 - 1.2 Simplest Strategy – Heuristic F
 - 1.3 Heuristic Board Evaluation Fur
 - 1.4 Rule-based Strategy
 - 1.5 Minimax Search Algorithm
 - 1.6 Minimax with Alpha-beta Prur



1. Playing Against Computer with AI (Advanced)

Click the image to run the demo for the various AI strategies (under the "Options" menu):



Tic-tac-toe seems dumb, but it actually requires you to *lookahead* one opponent's move to ensure that you will not lose. That is, you need to consider your opponent's move after your next move.

For example, suppose that the computer uses 'O'. At (D), the computer did not consider the opponent's next move and place at the corner (which is preferred over the side). At (E), the opponent was able to block the computer's winning move and create a fork.

x	x	x	x	x x
-----	-----	-----	-----	-----
	o	o	o	o
-----	-----	-----	-----	-----
		x	o x	o x
(A)	(B)	(C)	(D)	(E)

1.1 Implementing the AI Player

To test the various AI strategies, an abstract superclass called `AIPlayer` is defined, which takes the `Board` as an argument in its constructor (because you need the board position to compute the next move). An abstract method called `move()` is defined, which shall be implemented in subclasses with the chosen strategy.

```

1  /**
2   * Abstract superclass for all AI players with different strategies.
3   * To construct an AI player:
4   * 1. Construct an instance (of its subclass) with the game Board
5   * 2. Call setSeed() to set the computer's seed
6   * 3. Call move() which returns the next move in an int[2] array of {row, col}.
7   *
8   * The implementation subclasses need to override abstract method move().
9   * They shall not modify Cell[][], i.e., no side effect expected.
10  * Assume that next move is available, i.e., not game-over yet.

```

```

11  */
12  public abstract class AIPlayer {
13      protected int ROWS = GameMain.ROWS; // number of rows
14      protected int COLS = GameMain.COLS; // number of columns
15
16      protected Cell[][] cells; // the board's ROWS-by-COLS array of Cells
17      protected Seed mySeed;    // computer's seed
18      protected Seed oppSeed;   // opponent's seed
19
20      /** Constructor with reference to game board */
21      public AIPlayer(Board board) {
22          cells = board.cells;
23      }
24
25      /** Set/change the seed used by computer and opponent */
26      public void setSeed(Seed seed) {
27          this.mySeed = seed;
28          oppSeed = (mySeed == Seed.CROSS) ? Seed.NOUGHT : Seed.CROSS;
29      }
30
31      /** Abstract method to get next move. Return int[2] of {row, col} */
32      abstract int[] move(); // to be implemented by subclasses
33  }

```

1.2 Simplest Strategy – Heuristic Preferences via Table Lookup

The simplest computer strategy is to place the seed on the first empty cell in this order: the center, one of the four corners, one of the four sides. This dumb strategy, of course, does not work. It merely gets you started programming the computer play.

For example,

```

1  /**
2   * Computer move based on simple table lookup of preferences
3   */
4  public class AIPlayerTableLookup extends AIPlayer {
5
6      // Moves {row, col} in order of preferences. {0, 0} at top-left corner
7      private int[][] preferredMoves = {
8          {1, 1}, {0, 0}, {0, 2}, {2, 0}, {2, 2},
9          {0, 1}, {1, 0}, {1, 2}, {2, 1}};
10
11      /** constructor */
12      public AIPlayerTableLookup(Board board) {
13          super(board);
14      }
15
16      /** Search for the first empty cell, according to the preferences
17       * Assume that next move is available, i.e., not gameover
18       * @return int[2] of {row, col}
19       */
20      @Override
21      public int[] move() {
22          for (int[] move : preferredMoves) {
23              if (cells[move[0]][move[1]].content == Seed.EMPTY) {
24                  return move;
25              }
26          }
27          assert false : "No empty cell?!";
28          return null;
29      }
30  }

```

1.3 Heuristic Board Evaluation Function

In this strategy, we need to formula a *heuristic evaluation function*, which returns a relative score, e.g., $+\infty$ for

computer-win, $-\infty$ for opponent-win, 0 for neutral, and a number in between to indicate the relative advantage of the computer vs. the opponent.

In Tic-Tac-Toe, a possible heuristic evaluation function for the current board position is:

- +100 for EACH 3-in-a-line for computer.
- +10 for EACH two-in-a-line (with a empty cell) for computer.
- +1 for EACH one-in-a-line (with two empty cells) for computer.
- Negative scores for opponent, i.e., -100, -10, -1 for EACH opponent's 3-in-a-line, 2-in-a-line and 1-in-a-line.
- 0 otherwise (empty lines or lines with both computer's and opponent's seeds).

For Tic-Tac-Toe, compute the scores for each of the 8 lines (3 rows, 3 columns and 2 diagonals) and obtain the sum.

For an Othello (Reversi), the heuristic evaluation function could be the difference of computer's seeds over opponent's seeds.

To implement this strategy, you need to compute the score for all the valid moves, and place the seed at the position with the highest score. This strategy does not work in Tic-Tac-Toe (and in most of the board game) because it does not *lookahead* opponent's next move.

1.4 Rule-based Strategy

For Tic-tac-toe, the rules, in the order of importance, are:

- Rule 1: If I have a winning move, take it.
- Rule 2: If the opponent has a winning move, block it.
- Rule 3: If I can create a fork (two winning ways) after this move, do it.
- Rule 4: Do not let the opponent creating a fork after my move. (Opponent may block your winning move and create a fork.)
- Rule 5: Place in the position such as I may win in the most number of possible ways.

Rule 1 and 2 can be programmed easily. Rule 3 is harder. Rule 4 is even harder because you need to lookahead one opponent move, after your move. For rule 5, you need to count the number of possible winning ways.

Rule-based strategy is only applicable for simple game such as Tic-tac-toe and Othello.

1.5 Minimax Search Algorithm

Reference: Wiki "Minimax".

First, decide on a *heuristic board evaluation function* (see above section).

For Tic-Tac-Toe, the function could be as simple as returning +1 if the computer wins, -1 if the player wins, or 0 otherwise. However, simple evaluation function may require deeper search.

A better evaluation function for Tic-Tac-Toe is:

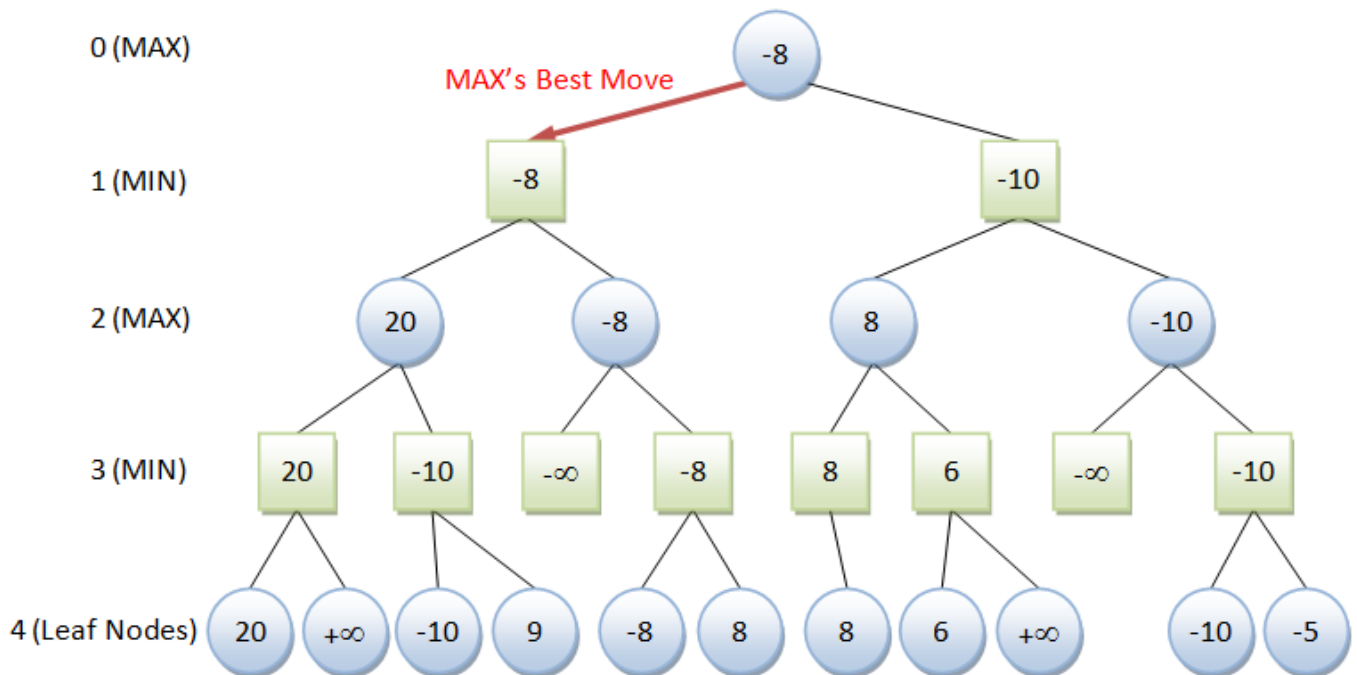
- +100 for EACH 3-in-a-line for computer.
- +10 for EACH 2-in-a-line (with a empty cell) for computer.
- +1 for EACH 1-in-a-line (with two empty cells) for computer.
- Negative scores for opponent, i.e., -100, -10, -1 for EACH opponent's 3-in-a-line, 2-in-a-line and 1-in-a-line.
- 0 otherwise (empty lines or lines with both computer's and opponent's seed).

Compute the scores for each of the 8 lines (3 rows, 3 columns and 2 diagonals) and obtain the sum.

The principle of minimax is to *minimize the maximum possible loss*.

As an illustration, suppose that there are only one or two possible moves per player in each turn. Furthermore, an evaluation function has been defined, which returns $+\infty$ if the computer wins, $-\infty$ if the computer loses, and a score in between to reflect the relative advantage of the computer. Computer (or the maximizing player) is represented by circle. The opponent (or the minimizing player) is represented by square. We limit the lookahead to four moves.

The algorithm evaluates the *leaf nodes* (terminating "gameover" nodes or at maximum depth of 4) using the heuristic evaluation function, obtaining the values shown. At level 3, the minimizing player will choose, for each node, the minimum of its children. In level 2, the maximizing player chooses the maximum of the children. The algorithm continues evaluating the maximum and minimum values of the child nodes alternately until it reaches the root node, where it chooses the move with the maximum value. This is the move that the player should make in order to minimize the maximum possible loss.



Minimax is *recursive*. The algorithm is as follows:

```

minimax(level, player) // player may be "computer" or "opponent"
if (gameover || level == 0)
    return score
children = all legal moves for this player
if (player is computer, i.e., maximizing player)
    // find max
    bestScore = -inf
    for each child
        score = minimax(level - 1, opponent)
        if (score > bestScore) bestScore = score
    return bestScore
else (player is opponent, i.e., minimizing player)
    // find min
    bestScore = +inf
    for each child
        score = minimax(level - 1, computer)
        if (score < bestScore) bestScore = score
    return bestScore

// Initial Call
minimax(2, computer)

```

```

1  import java.util.*;
2  /** AIPlayer using Minimax algorithm */
3  public class AIPlayerMinimax extends AIPlayer {
4
5      /** Constructor with the given game board */
6      public AIPlayerMinimax(Board board) {
7          super(board);
8      }
9
10     /** Get next best move for computer. Return int[2] of {row, col} */
11     @Override
12     int[] move() {

```

```

13     int[] result = minimax(2, mySeed); // depth, max turn
14     return new int[] {result[1], result[2]}; // row, col
15 }
16
17 /** Recursive minimax at level of depth for either maximizing or minimizing player.
18     Return int[3] of {score, row, col} */
19 private int[] minimax(int depth, Seed player) {
20     // Generate possible next moves in a List of int[2] of {row, col}.
21     List<int[]> nextMoves = generateMoves();
22
23     // mySeed is maximizing; while oppSeed is minimizing
24     int bestScore = (player == mySeed) ? Integer.MIN_VALUE : Integer.MAX_VALUE;
25     int currentScore;
26     int bestRow = -1;
27     int bestCol = -1;
28
29     if (nextMoves.isEmpty() || depth == 0) {
30         // Gameover or depth reached, evaluate score
31         bestScore = evaluate();
32     } else {
33         for (int[] move : nextMoves) {
34             // Try this move for the current "player"
35             cells[move[0]][move[1]].content = player;
36             if (player == mySeed) { // mySeed (computer) is maximizing player
37                 currentScore = minimax(depth - 1, oppSeed)[0];
38                 if (currentScore > bestScore) {
39                     bestScore = currentScore;
40                     bestRow = move[0];
41                     bestCol = move[1];
42                 }
43             } else { // oppSeed is minimizing player
44                 currentScore = minimax(depth - 1, mySeed)[0];
45                 if (currentScore < bestScore) {
46                     bestScore = currentScore;
47                     bestRow = move[0];
48                     bestCol = move[1];
49                 }
50             }
51             // Undo move
52             cells[move[0]][move[1]].content = Seed.EMPTY;
53         }
54     }
55     return new int[] {bestScore, bestRow, bestCol};
56 }
57
58 /** Find all valid next moves.
59     Return List of moves in int[2] of {row, col} or empty list if gameover */
60 private List<int[]> generateMoves() {
61     List<int[]> nextMoves = new ArrayList<int[]>(); // allocate List
62
63     // If gameover, i.e., no next move
64     if (hasWon(mySeed) || hasWon(oppSeed)) {
65         return nextMoves; // return empty list
66     }
67
68     // Search for empty cells and add to the List
69     for (int row = 0; row < ROWS; ++row) {
70         for (int col = 0; col < COLS; ++col) {
71             if (cells[row][col].content == Seed.EMPTY) {
72                 nextMoves.add(new int[] {row, col});
73             }
74         }
75     }
76     return nextMoves;
77 }
78
79 /** The heuristic evaluation function for the current board
80     @Return +100, +10, +1 for EACH 3-, 2-, 1-in-a-line for computer.
81     -100, -10, -1 for EACH 3-, 2-, 1-in-a-line for opponent.

```

```

82         0 otherwise */
83 private int evaluate() {
84     int score = 0;
85     // Evaluate score for each of the 8 lines (3 rows, 3 columns, 2 diagonals)
86     score += evaluateLine(0, 0, 0, 1, 0, 2); // row 0
87     score += evaluateLine(1, 0, 1, 1, 1, 2); // row 1
88     score += evaluateLine(2, 0, 2, 1, 2, 2); // row 2
89     score += evaluateLine(0, 0, 1, 0, 2, 0); // col 0
90     score += evaluateLine(0, 1, 1, 1, 2, 1); // col 1
91     score += evaluateLine(0, 2, 1, 2, 2, 2); // col 2
92     score += evaluateLine(0, 0, 1, 1, 2, 2); // diagonal
93     score += evaluateLine(0, 2, 1, 1, 2, 0); // alternate diagonal
94     return score;
95 }
96
97 /** The heuristic evaluation function for the given line of 3 cells
98     @Return +100, +10, +1 for 3-, 2-, 1-in-a-line for computer.
99     -100, -10, -1 for 3-, 2-, 1-in-a-line for opponent.
100    0 otherwise */
101 private int evaluateLine(int row1, int col1, int row2, int col2, int row3, int col3) {
102     int score = 0;
103
104     // First cell
105     if (cells[row1][col1].content == mySeed) {
106         score = 1;
107     } else if (cells[row1][col1].content == oppSeed) {
108         score = -1;
109     }
110
111     // Second cell
112     if (cells[row2][col2].content == mySeed) {
113         if (score == 1) { // cell1 is mySeed
114             score = 10;
115         } else if (score == -1) { // cell1 is oppSeed
116             return 0;
117         } else { // cell1 is empty
118             score = 1;
119         }
120     } else if (cells[row2][col2].content == oppSeed) {
121         if (score == -1) { // cell1 is oppSeed
122             score = -10;
123         } else if (score == 1) { // cell1 is mySeed
124             return 0;
125         } else { // cell1 is empty
126             score = -1;
127         }
128     }
129
130     // Third cell
131     if (cells[row3][col3].content == mySeed) {
132         if (score > 0) { // cell1 and/or cell2 is mySeed
133             score *= 10;
134         } else if (score < 0) { // cell1 and/or cell2 is oppSeed
135             return 0;
136         } else { // cell1 and cell2 are empty
137             score = 1;
138         }
139     } else if (cells[row3][col3].content == oppSeed) {
140         if (score < 0) { // cell1 and/or cell2 is oppSeed
141             score *= 10;
142         } else if (score > 1) { // cell1 and/or cell2 is mySeed
143             return 0;
144         } else { // cell1 and cell2 are empty
145             score = -1;
146         }
147     }
148     return score;
149 }
150

```

```

151     private int[] winningPatterns = {
152         0b111000000, 0b000111000, 0b000000011, // rows
153         0b100100100, 0b010010010, 0b001001001, // cols
154         0b100010001, 0b001010100                // diagonals
155     };
156
157     /** Returns true if thePlayer wins */
158     private boolean hasWon(Seed thePlayer) {
159         int pattern = 0b000000000; // 9-bit pattern for the 9 cells
160         for (int row = 0; row < ROWS; ++row) {
161             for (int col = 0; col < COLS; ++col) {
162                 if (cells[row][col].content == thePlayer) {
163                     pattern |= (1 << (row * COLS + col));
164                 }
165             }
166         }
167         for (int winningPattern : winningPatterns) {
168             if ((pattern & winningPattern) == winningPattern) return true;
169         }
170         return false;
171     }
172 }


```

Note: The pseudocode presented in Wiki "minimax" is known as "negamax", which is very hard to understand, and even harder to program and debug.

1.6 Minimax with Alpha-beta Pruning

Reference: Wiki "Alpha-beta pruning".

Alpha-beta pruning seeks to reduce the number of nodes that needs to be evaluated in the search tree by the minimax algorithm. For example in the alpha cut-off, since node D returns 1, node C (MIN) cannot be more than 1. But node B is 4. There is no need to search the other children of node C, as node A will certainly pick node B over node C.



In the algorithm, two parameters are needed: an alpha value which holds the best MAX value found for MAX node; and a beta value which holds the best MIN value found for MIN node. As illustrated, the remaining children can be aborted if $\alpha \geq \beta$, for both the alpha cut-off and beta cut-off. Alpha and beta are initialized to $-\infty$ and $+\infty$ respectively.

The recursive algorithm for "minimax with alpha-beta pruning" is as follows:

```

minimax(level, player, alpha, beta) // player may be "computer" or "opponent"
if (gameover || level == 0)
    return score
children = all valid moves for this "player"
if (player is computer, i.e., max's turn)
    // Find max and store in alpha
    for each child
        score = minimax(level - 1, opponent, alpha, beta)
        if (score > alpha) alpha = score
        if (alpha >= beta) break; // beta cut-off
    return alpha
else (player is opponent, i.e., min's turn)

```

```

// Find min and store in beta
for each child
    score = minimax(level - 1, computer, alpha, beta)
    if (score < beta) beta = score
    if (alpha >= beta) break; // alpha cut-off
return beta

// Initial call with alpha=-inf and beta=inf
minimax(2, computer, -inf, +inf)

```

The relevant changes (over the AIPlayerMinimax.java) are:

```

/** Get next best move for computer. Return int[2] of {row, col} */
@Override
int[] move() {
    int[] result = minimax(2, mySeed, Integer.MIN_VALUE, Integer.MAX_VALUE);
    // depth, max-turn, alpha, beta
    return new int[] {result[1], result[2]}; // row, col
}

/** Minimax (recursive) at level of depth for maximizing or minimizing player
    with alpha-beta cut-off. Return int[3] of {score, row, col} */
private int[] minimax(int depth, Seed player, int alpha, int beta) {
    // Generate possible next moves in a list of int[2] of {row, col}.
    List<int[]> nextMoves = generateMoves();

    // mySeed is maximizing; while oppSeed is minimizing
    int score;
    int bestRow = -1;
    int bestCol = -1;

    if (nextMoves.isEmpty() || depth == 0) {
        // Gameover or depth reached, evaluate score
        score = evaluate();
        return new int[] {score, bestRow, bestCol};
    } else {
        for (int[] move : nextMoves) {
            // try this move for the current "player"
            cells[move[0]][move[1]].content = player;
            if (player == mySeed) { // mySeed (computer) is maximizing player
                score = minimax(depth - 1, oppSeed, alpha, beta)[0];
                if (score > alpha) {
                    alpha = score;
                    bestRow = move[0];
                    bestCol = move[1];
                }
            } else { // oppSeed is minimizing player
                score = minimax(depth - 1, mySeed, alpha, beta)[0];
                if (score < beta) {
                    beta = score;
                    bestRow = move[0];
                    bestCol = move[1];
                }
            }
            // undo move
            cells[move[0]][move[1]].content = Seed.EMPTY;
            // cut-off
            if (alpha >= beta) break;
        }
        return new int[] {(player == mySeed) ? alpha : beta, bestRow, bestCol};
    }
}

```

REFERENCES & RESOURCES

1. Wiki "Minimax" and "Alpha-beta pruning"
2. Dong Xiang, "Solve Tic Tac Toe with the MiniMax algorithm" @ <http://www.codeproject.com/Articles/43622/Solve-Tic-Tac-Toe-with-the-MiniMax-algorithm>.
3. "Minimax Explained" @ <http://ai-depot.com/articles/minimax-explained>.

Latest version tested: JDK 1.7.2
Last modified: May, 2012

Feedback, comments, corrections, and errata can be sent to Chua Hock-Chuan (ehchua@ntu.edu.sg) | [HOME](#)